

Data Wrangling

Sofia

```
# load packages
library(tidyverse)
library(kableExtra)
library(robotstxt)
library(rvest)
library(purrr)
library(readr)
library(tidyr)
library(sf) #for reading shape files

# set code chunk defaults
knitr::opts_chunk$set(tidy = F, # display code as typed
                      size = "small", # slightly smaller code font
                      message = FALSE,
                      warning = FALSE,
                      comment = "\t")

# set black & white default plot theme
theme_set(theme_classic())

# improve digit and NA display
options(scipen = 1, knitr.kable.NA = '')
```

Data Source 1: Kaggle- Fuel Data

Cleaning

```
'data.frame':  38113 obs. of   5 variables:
```

\$ year	:	int	1984 1984 1984 1984 1984 1984 1984 1984 1984 1984 1984 ...
\$ engine_cylinders	:	int	6 6 4 4 4 4 6 6 6 6 ...
\$ fuel_type	:	chr	"Regular" "Regular" "Regular" "Regular" ...
\$ city_mpg_ft1	:	int	17 17 18 18 18 18 13 13 15 15 ...
\$ tailpipe_co2_in_grams_mile_ft1	:	num	444 444 423 423 523 ...

```
# Select only the specified columns
Fuel_hclust <- Fuel_clean[, c("year", "engine_cylinders", "city_mpg_ft1", "tailpipe_co2_in

# View the structure of the cleaned dataset
str(Fuel_hclust)
```

2


```

[445] 1 2 2 1 1 1 2 2 1 2 1 1 1 1 2 2 1 2 2 2 1 2 2 1 1 1 3 1 1 2 1 1 1 1 2 1 1
[482] 1 1 1 1 1 1 2 1 2 2 1 2 1 1 2 1 1 1 1 1 2 1 2 1 1 1 1 1 2 1 1 1 1 1 1 1
[519] 2 2 1 1 1 1 1 1 1 1 1 2 1 2 1 1 2 1 1 2 1 1 1 3 1 1 1 2 1 1 1 1 1 2
[556] 2 1 2 2 2 1 2 2 1 3 1 1 2 1 1 2 1 1 2 1 1 1 1 1 1 1 2 1 1 1 1 1 2 1
[593] 1 1 1 2 2 1 1 1 1 1 1 1 1 1 2 2 2 3 2 1 2 1 1 2 2 1 1 2 1 1 1 2 1 2 1 1 2
[630] 2 2 2 1 1 2 2 1 2 2 2 1 1 1 1 1 2 2 1 1 1 1 1 1 1 1 2 2 1 3 1 1 1 1 2 2
[667] 1 1 1 1 1 1 2 2 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 2 1 1 1 1 2 1 1 1 1 2 2
[704] 2 2 3 1 1 1 1 1 2 1 1 2 2 1 2 1 1 1 1 1 2 1 1 3 1 2 1 2 2 2 1 1 1 1 2 1 1
[741] 1 1 2 1 2 2 1 1 1 1 2 1 1 1 1 1 1 1 1 2 1 1 1 1 2 1 1 2 1 2 1 3 2 2 1 2 1
[778] 1 2 1 1 2 2 2 1 1 2 1 1 2 2 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 2 2 1 2 2 2 1
[815] 2 1 1 1 1 2 2 2 1 2 1 2 1 1 1 1 2 2 2 1 1 2 1 1 1 1 2 1 1 2 2 1 1 1 1 1 1
[852] 1 1 2 1 1 1 2 1 1 3 2 1 1 1 1 2 2 2 2 1 3 1 2 1 1 2 1 1 1 1 2 1 1 3 1 1 1
[889] 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 2 2 1 2 1 1 1 1 1 2 2 2 1 2 1 1 1
[926] 1 1 1 1 1 1 1 2 1 1 1 2 1 1 2 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 3 1 2 1 1 2
[963] 1 1 1 1 1 1 2 1 1 1 1 2 1 2 1 1 1 1 2 2 1 1 1 1 1 1 2 1 1 3 1 1 1 1 2 1 1
[1000] 1

```

```

# Perform hierarchical clustering with complete linkage
hclust_result <- hclust(dist(scaled_data), method = "complete")

# View cluster assignments
cluster_assignments <- cutree(hclust_result, k = 3) # Adjust k as needed
print(cluster_assignments)

```

```

[1] 1 1 1 1 2 2 2 2 1 1 1 3 1 1 1 1 1 1 2 2 1 1 2 2 1 1 1 1 1 1 1 1 1 1 2 1 1
[38] 1 1 1 2 1 2 2 1 1 1 1 1 2 1 2 2 1 1 1 1 1 1 2 2 1 1 1 1 1 1 1 2 1 1 1 1 1
[75] 2 1 1 1 1 3 1 1 2 1 1 2 2 1 1 2 1 1 1 2 1 1 1 3 2 1 2 2 1 1 1 1 1 2 1 2
[112] 1 1 1 1 1 1 1 2 1 2 2 3 1 1 1 1 1 1 1 1 1 2 1 2 1 2 1 2 2 1 1 2 1 1 1 2
[149] 1 1 1 2 1 2 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 2 2 2 1 1 1 1 1 1 2
[186] 1 2 1 1 1 2 2 1 2 1 1 1 1 1 1 1 1 1 2 2 1 3 2 1 2 1 2 2 1 1 1 2 1 1 2 1 2
[223] 2 1 2 1 2 2 1 1 1 2 2 2 2 1 2 1 2 1 1 1 1 2 1 1 2 1 1 1 2 1 1 1 1 2 2 2 2
[260] 1 2 1 2 1 1 2 2 1 3 1 1 1 2 1 1 2 2 1 1 1 1 1 1 1 1 1 2 2 1 1 1 2 1 1 2 1 1
[297] 1 1 2 3 1 2 3 1 1 1 1 1 1 1 2 1 2 2 1 2 1 2 1 1 1 1 1 1 1 1 1 2 1 1 2 1 1 1
[334] 1 2 2 2 1 1 1 1 2 1 1 1 1 1 2 2 2 1 2 1 1 1 1 1 1 1 3 1 2 1 2 2 1 1 2 1 1
[371] 1 1 2 2 1 1 3 1 1 1 1 1 1 1 1 2 2 2 2 2 1 1 1 2 1 1 2 2 1 1 2 1 2 2 1 1 1
[408] 2 1 1 2 1 1 2 1 1 1 2 1 1 2 1 2 1 2 1 1 1 2 1 2 1 2 2 1 1 1 2 1 2 2 1 1 1
[445] 1 2 2 1 1 1 2 2 1 2 1 1 1 1 2 2 1 2 2 2 1 2 2 1 1 1 3 1 1 2 1 1 1 1 2 1 1
[482] 1 1 1 1 1 1 2 1 2 2 1 2 1 1 2 1 1 1 1 1 1 2 1 2 1 1 1 1 1 2 1 1 1 1 1 1 1
[519] 2 2 1 1 1 1 1 1 1 1 1 2 1 2 1 1 2 1 1 2 1 1 1 3 1 1 1 2 1 1 1 1 1 1 2
[556] 2 1 2 2 2 1 2 2 1 3 1 1 2 1 1 2 1 1 2 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 2 1
[593] 1 1 1 2 2 1 1 1 1 1 1 1 1 1 2 2 2 3 2 1 2 1 1 2 2 1 1 2 1 1 1 2 1 2 1 1 2

```

```

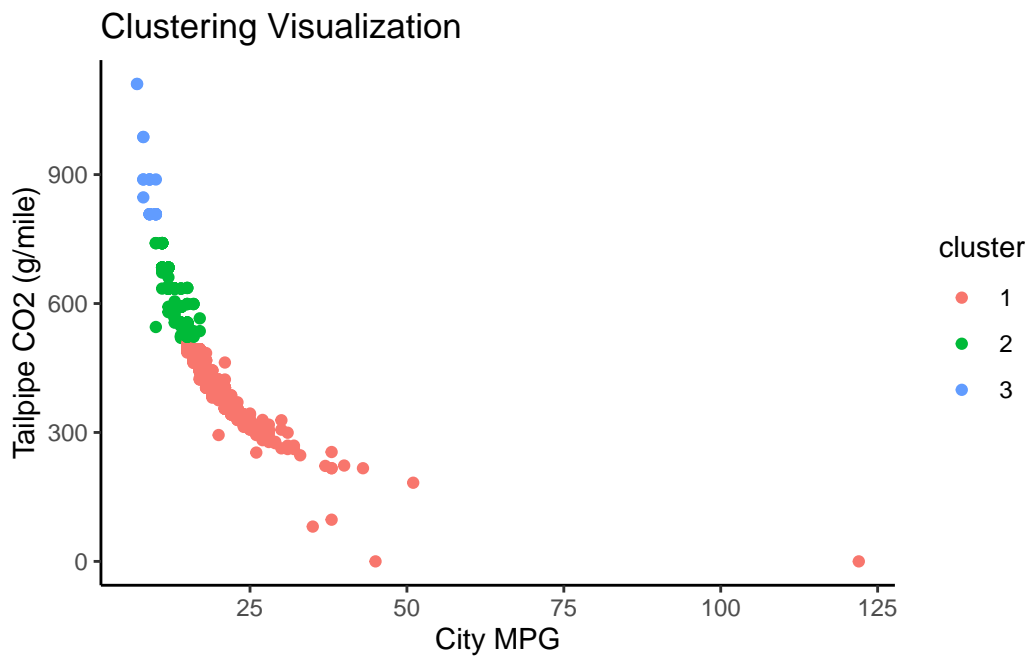
[630] 2 2 2 1 1 2 2 1 2 2 2 1 1 1 1 1 2 2 1 1 1 1 1 1 1 2 2 1 3 1 1 1 1 1 2 2
[667] 1 1 1 1 1 1 2 2 1 1 1 1 1 1 1 2 1 1 1 1 1 2 1 1 1 2 1 1 2 1 1 1 1 2 2
[704] 2 2 3 1 1 1 1 1 2 1 1 2 2 1 2 1 1 1 1 1 2 1 1 3 1 2 1 2 2 2 1 1 1 1 2 1 1
[741] 1 1 2 1 2 2 1 1 1 1 2 1 1 1 1 1 1 1 1 2 1 1 1 1 2 1 1 2 1 2 1 3 2 2 1 2 1
[778] 1 2 1 1 2 2 2 1 1 2 1 1 2 2 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 2 2 1 2 2 2 1
[815] 2 1 1 1 1 2 2 2 1 2 1 2 1 1 1 1 2 2 2 1 1 2 1 1 1 1 2 1 1 2 2 1 1 1 1 1
[852] 1 1 2 1 1 1 2 1 1 3 2 1 1 1 1 2 2 2 2 1 3 1 2 1 1 2 1 1 1 1 2 1 1 3 1 1 1
[889] 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 2 2 1 2 1 1 1 1 1 2 2 2 1 2 1 1 1
[926] 1 1 1 1 1 1 1 2 1 1 1 2 1 1 2 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 3 1 2 1 1 2
[963] 1 1 1 1 1 1 2 1 1 1 1 2 1 2 1 1 1 1 2 2 1 1 1 1 1 1 2 1 1 3 1 1 1 1 2 1 1
[1000] 1

```

```
library(ggplot2)
```


```
# Add cluster assignments to the sampled data
sampled_data$cluster <- factor(cluster_assignments)
```

```
# Plot the data points with cluster assignments
ggplot(sampled_data, aes(x = city_mpg_ft1, y = tailpipe_co2_in_grams_mile_ft1, color = cluster)) +
  geom_point() +
  labs(title = "Clustering Visualization", x = "City MPG", y = "Tailpipe CO2 (g/mile)")
```




```
# Plot the dendrogram
plot(hclust_result)
```

Cluster Dendrogram



```
dist(scaled_data)
hclust (*, "complete")
```

```
suppressPackageStartupMessages(library(dendextend))
avg_dend_obj <- as.dendrogram(hclust_result)
avg_col_dend <- color_branches(avg_dend_obj, h = 3)
plot(avg_col_dend)
```



```
# Load the required libraries
library(readr)
```

```

library(dplyr)
library(ggplot2)

# Select relevant variable
variable <- "tailpipe_co2_in_grams_mile_ft1"

# Sample the dataset (optional)
set.seed(123) # Set seed for reproducibility
sampled_data <- Fuel_hclust %>% sample_n(500) # Adjust the number of samples as needed

# Normalize the data (optional)
scaled_data <- scale(sampled_data[[variable]])

# Compute the distance matrix
distance_matrix <- dist(scaled_data)

# Perform hierarchical clustering
hierarchical_clusters <- hclust(distance_matrix, method = "ward.D2")

# Plot the dendrogram
plot(hierarchical_clusters, main = "Dendrogram of Hierarchical Clustering", xlab = "", ylab = "Distance")

```

Dendrogram of Hierarchical Clustering



`hclust (*, "ward.D2")`

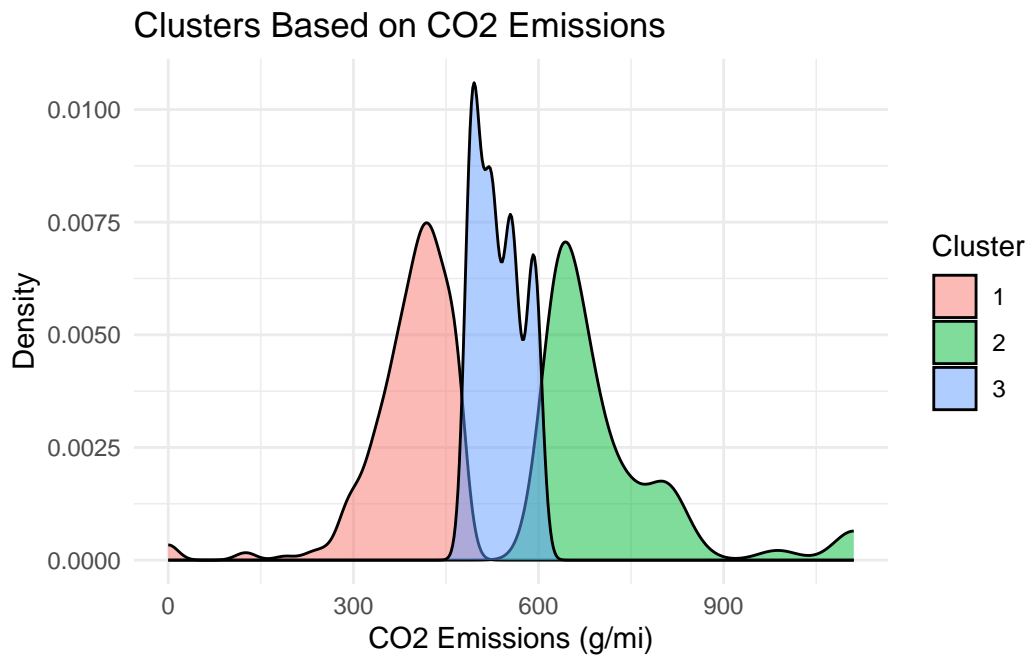
```

# Cut the dendrogram to get clusters
num_clusters <- 3 # You can adjust this based on the dendrogram
clusters <- cutree(hierarchical_clusters, k = num_clusters)

# Add cluster labels to the original dataset
sampled_data$cluster <- clusters

# Visualize the clusters
ggplot(sampled_data, aes(x = tailpipe_co2_in_grams_mile_ft1, y = ..density.., fill = factor(cluster))) +
  geom_density(alpha = 0.5) +
  labs(x = "CO2 Emissions (g/mi)", y = "Density", title = "Clusters Based on CO2 Emissions") +
  scale_fill_discrete(name = "Cluster") +
  theme_minimal()

```



```
library(dplyr)
```