

# Prep2S24

Sofia Guttmann

2024-02-08

Reminder: Prep assignments are to be completed individually. Upload a final copy of the .Qmd and renamed .pdf to your private repo, and submit the renamed pdf to Gradescope before Sunday, Feb. 11th at midnight (11:59 pm is what Gradescope shows).

## Reading

The associated reading for the week is Chapter 4, Chapter 5, Chapter 6 (skip 6.4) and Sections 8.3 and 8.4. This reading explores major functions in wrangling data, including reshaping data. There are many commands here to learn about - do your best to develop a sense of what they each do, and we will build on that by using them for the rest of the semester. You don't need to memorize them all.

Remember, I recommend you code along with the book examples. You can try out the code yourself - just be sure to load the mdsr package and any other packages referenced. You can get the code in R script files (basically, files of just R code, not like a .Rmd or .Qmd) from the book website.

## 1 - Some basics

Many different data wrangling commands are covered in these chapters. Identify the command you'd use for each of the operations below.

part a - Add the average of 3 variables to the data set as a new variable.

Solution: You can use the **mutate()** function from the **dplyr** package

part b - Keep only 4 columns of a data frame in a new data set.

Solution: You can use the **select()** function from the **dplyr** package

part c - Choose observations that match a particular category of a categorical variable to keep in a new data set.

Solution: You can use the **filter()** function from the **dplyr** package

part d - Combine two data sets based on common variables where all rows from the first are returned, along with any matches in the second.

Solution: You can use the **left\_join()** function from the **dplyr** package

## 2 - NYC Flights

In Section 5.1, the `flights` and `airlines` tables within the `nycflights13` package are joined together.

part a - Recreate the `flights_joined` dataset from Section 5.1, being sure to *glimpse* the data in the Console (or via the code chunk) to verify the join worked.

Solution:

```
library(dplyr)
library(nycflights13)

flights_joined <- inner_join(flights, airlines, by = "carrier")

print(head(flights_joined))
```

```
# A tibble: 6 x 20
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517           515           2     830           819
2  2013     1     1     533           529           4     850           830
3  2013     1     1     542           540           2     923           850
4  2013     1     1     544           545          -1    1004          1022
5  2013     1     1     554           600          -6     812           837
6  2013     1     1     554           558          -4     740           728
# i 12 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>, name <chr>
```

part b - Now, starting from `flights_joined`, create a new dataset `flights_short` that does the following:

- creates a new variable, `distance_km`, which is distance in kilometers (note that 1 mile is about 1.6 kilometers)
- keeps only the variables: `name`, `flight`, `arr_delay`, and `distance_km` and
- keeps only observations where the distance is less than 480 kilometers (300 miles).

Solution:

```
flights_joined <- flights_joined %>%
  mutate(distance_km = distance * 1.6)
```

```

flights_short <- flights_joined %>%
  select(name, flight, arr_delay, distance_km) %>%
  filter(distance_km < 480)

print(head(flights_short))

```

```

# A tibble: 6 x 4
  name                flight arr_delay distance_km
<chr>                <int>    <dbl>      <dbl>
1 ExpressJet Airlines Inc.  5708      -14        366.
2 JetBlue Airways         1806       -4        299.
3 Southwest Airlines Co.   4646     -19        296
4 ExpressJet Airlines Inc.  4144       12        339.
5 JetBlue Airways         1002     -10        299.
6 JetBlue Airways          20       -1        422.

```

part c - Using the functions introduced in Section 4.1.4, compute the number of flights (call this `N`), the average arrival delay (call this `avg_arr_delay`), and the average distance in kilometers (call this `avg_dist_km`) among these flights with distances less than 480 km (i.e. working off of `flights_short`), grouping by the carrier name. Sort the results in descending order based on `avg_arr_delay`. Save the results in a tibble object called `delay_summary`, and display the table.

Solution:

```

library(dplyr)
library(tibble)

delay_summary <- flights_short %>%
  group_by(name) %>%
  summarize(N = n(),
            avg_arr_delay = mean(arr_delay, na.rm = TRUE),
            avg_dist_km = mean(distance_km, na.rm = TRUE)) %>%
  arrange(desc(avg_arr_delay))

print(delay_summary)

```

```

# A tibble: 11 x 4
  name                N avg_arr_delay avg_dist_km

```

	<chr>	<int>	<dbl>	<dbl>
1	Mesa Airlines Inc.	319	18.0	361.
2	ExpressJet Airlines Inc.	15588	15.6	372.
3	Envoy Air	2924	11.0	350.
4	JetBlue Airways	10813	8.36	360.
5	Endeavor Air Inc.	5779	6.84	319.
6	Southwest Airlines Co.	208	4.92	272.
7	United Air Lines Inc.	3353	4.09	320.
8	SkyWest Airlines Inc.	1	3	366.
9	US Airways Inc.	9633	2.22	309.
10	American Airlines Inc.	1455	1.88	299.
11	Delta Air Lines Inc.	1214	-0.643	325.

part d - Rename the four columns in the `delay_summary` data table to `Airline`, `"Total flights under 480 km"`, `"Average arrival delay (mins)"` and `"Average distance (km)"`, respectively, then use `kable(booktabs = TRUE, digits = 0)` to make the final table output in the pdf close to publication quality.

Solution:

```
library(knitr)

names(delay_summary) <- c("Airline", "Total flights under 480 km", "Average arrival delay", "Average distance (km)")

kable(delay_summary, booktabs = TRUE, digits = 0)
```

Airline	Total flights under 480 km	Average arrival delay (mins)	Average distance (km)
Mesa Airlines Inc.	319	18	361
ExpressJet Airlines Inc.	15588	16	372
Envoy Air	2924	11	350
JetBlue Airways	10813	8	360
Endeavor Air Inc.	5779	7	319
Southwest Airlines Co.	208	5	272
United Air Lines Inc.	3353	4	320
SkyWest Airlines Inc.	1	3	366
US Airways Inc.	9633	2	309
American Airlines Inc.	1455	2	299
Delta Air Lines Inc.	1214	-1	325

### 3 - Baby names - Variant of 6.2.5 example

part a - Working with the `babynames` data in the `babynames` package, create a dataset `recent_names` that only includes years 2003 to 2017 (giving us the most recent 15 years of data).

Solution:

```
library(babynames)
library(dplyr)

recent_names <- babynames %>%
  filter(year >= 2003 & year <= 2017)

print(head(recent_names))
```

```
# A tibble: 6 x 5
  year sex  name      n  prop
<dbl> <chr> <chr>  <int> <dbl>
1  2003 F    Emily  25688 0.0128
2  2003 F    Emma   22704 0.0113
3  2003 F   Madison 20197 0.0101
4  2003 F   Hannah 17634 0.00879
5  2003 F   Olivia 16146 0.00805
6  2003 F  Abigail 15925 0.00794
```

part b - Following the code presented in Section 6.2.5, create a dataset called `recentnames_summary` that summarizes the total number of people in recent history (years 2003 to 2017) with each name, grouped by sex.

Solution:

```
library(babynames)
library(dplyr)

recent_names <- babynames %>%
  filter(year >= 2003 & year <= 2017)

recentnames_summary <- recent_names %>%
  group_by(name, sex) %>%
```

```
summarize(total_people = sum(n, na.rm = TRUE))
```

`summarise()` has grouped output by 'name'. You can override using the `.groups` argument.

```
print(head(recentnames_summary))
```

```
# A tibble: 6 x 3
# Groups:   name [6]
  name      sex  total_people
  <chr>    <chr>         <int>
1 Aaban    M             107
2 Aabha    F              35
3 Aabid    M              10
4 Aabir    M               5
5 Aabriella F             32
6 Aada     F               5
```

part c - Now, following the third and fourth code chunks presented in Section 6.2.5, reshape or *pivot* the summary data from *long* format to *wide* format. Only keep observations where more than 8,000 babies have been named in each sex (M and F), and find the smaller of the two ratios M / F and F / M to identify the top three sex-balanced names (and only the top three!). Save the wide data as `recentnames_balanced_wide`. Display the table.

Solution:

```
recentnames_summary_wide <- recentnames_summary %>%
  pivot_wider(names_from = sex, values_from = total_people)

recentnames_summary_wide_filtered <- recentnames_summary_wide %>%
  filter(M > 8000 & F > 8000)

recentnames_summary_wide_filtered <- recentnames_summary_wide_filtered %>%
  mutate(ratio_M_F = M / F,
         ratio_F_M = F / M)
names
```

```
function (x) .Primitive("names")
```

```

top_three_balanced_names <- recentnames_summary_wide_filtered %>%
  mutate(min_ratio = pmin(ratio_M_F, ratio_F_M)) %>%
  top_n(3, min_ratio)

recentnames_balanced_wide <- top_three_balanced_names

print(recentnames_balanced_wide)

```

```

# A tibble: 26 x 6
# Groups:   name [26]
   name      M      F ratio_M_F ratio_F_M min_ratio
   <chr> <int> <int>    <dbl>    <dbl>    <dbl>
1 Alexis 29951 118367  0.253     3.95     0.253
2 Amari  14612  9983   1.46     0.683    0.683
3 Angel 125664 26344   4.77     0.210    0.210
4 Avery  28018 102487  0.273     3.66     0.273
5 Cameron 116636 12529   9.31     0.107    0.107
6 Casey  12033  8278   1.45     0.688    0.688
7 Charlie 19563 12818   1.53     0.655    0.655
8 Dakota 23823 18971   1.26     0.796    0.796
9 Dylan  173360 9267  18.7     0.0535   0.0535
10 Emerson 11125 17993   0.618     1.62     0.618
# i 16 more rows

```

part d - Finally, use `pivot_longer()` to put the dataset back into *long* form. Call this dataset `recentnames_balanced` and display the table.

Solution:

```

recentnames_balanced <- recentnames_balanced_wide %>%
  pivot_longer(cols = c(M, F, ratio_M_F, ratio_F_M, min_ratio),
    names_to = "Variable",
    values_to = "Value")

print(recentnames_balanced)

```

```

# A tibble: 130 x 3
# Groups:   name [26]
   name Variable      Value

```



	<chr>	<chr>	<dbl>
1	Alexis	M	29951
2	Alexis	F	118367
3	Alexis	ratio_M_F	0.253
4	Alexis	ratio_F_M	3.95
5	Alexis	min_ratio	0.253
6	Amari	M	14612
7	Amari	F	9983
8	Amari	ratio_M_F	1.46
9	Amari	ratio_F_M	0.683
10	Amari	min_ratio	0.683

# i 120 more rows

## 4 - Ethics

Each subsection of Section 8.4 discusses an ethical scenario and ends with one or more questions. Consider the subsection 8.4.6 on “Reproducible spreadsheet analysis”.

Write two or three sentences reflecting on how using RMarkdown would help avoid some of the issues described in this scenario, or at least make them easier to spot.

Solution: Using RMarkdown allows for reproducibility and transparency. Using one file ensures ease when tracking changes and transformations.